




hugs id



— TECHNICAL REFERENCE · V 1 . 0

HUGS ID™


Technical Specification

The Canonical Identity and Governance Framework for the AI Era

Version 1.0 · Public Working Draft

4 August 2026 — Publisher: HUGS ID™

Canonical reference: hugsid.com · Documentation portal: learn.xpguess.com



Status of This Document

This document is the authoritative technical reference for the HUGS ID™ platform. It defines the architecture, identity model, governance framework, consent framework, AI permission model, trust framework, trust chain, and application programming interfaces of the HUGS ID standard.

Document	HUGS ID™ Technical Specification
Version	1.0 (semantic versioning applies to all schemas defined herein)
Status	Public Working Draft — open for review and implementation feedback
Publication date	4 August 2026
Publisher	HUGS ID™
Canonical reference	hugsid.com
Documentation portal	learn.xpguess.com
Intended audience	Developers, researchers, institutions, governing bodies, AI systems

This is a working draft. It does not yet constitute a final standard. Implementers should expect editorial refinement, while the core object models defined in Sections 4 through 9 are considered stable within the v1.x line. Comments and implementation reports are welcomed through the documentation portal and will inform subsequent public drafts.

Document Conventions

The key words MUST, MUST NOT, REQUIRED, SHALL, SHALL NOT, SHOULD, SHOULD NOT, RECOMMENDED, MAY, and OPTIONAL in this document are to be interpreted as described in RFC 2119 [1]. Requirement identifiers take the form REQ-<section>-<number> and are stable across v1.x revisions. All timestamps are ISO 8601 / RFC 3339 UTC. All object schemas are version-controlled and published under the reserved namespace hugsid.com/schemas/v1.

Object names and field names appear in monospace-style lowercase snake_case when referenced in prose. Figures and tables are numbered sequentially within each section.

Table of Contents

1. Vision.....	3
1.1 Vision Statement	4
1.2 Objectives	5

1.3 Non-Goals	6
2. Problem Statement	7
2.1 The Fragmentation of Digital Identity.....	8
2.2 The AI Inflection Point	9
2.3 Design Principles	10
3. Architecture	11
3.1 Layered Architecture	12
3.2 Layer Responsibilities	13
3.3 Data Flow	14
3.4 Governed Domains.....	15
4. Identity Layer.....	16
4.1 The Canonical Identity Object	17
4.2 Identifier Scheme	18
4.3 Identity Types	19
4.4 Identity Lifecycle.....	20
5. Governance Layer	21
5.1 Governance Objects	22
5.2 Policy Model	23
5.3 Authority Chains	24
5.4 Institutional Roles.....	25
6. Consent Layer	26
6.1 The Consent Object	27
6.2 Purpose Taxonomy	28
6.3 Consent Lifecycle	29
6.4 Revocation	30
7. Athlete Intelligence	31
7.1 Overview	32
7.2 Governed Data Classes	33
7.3 Use Cases.....	34
8. AI Permissions	35
8.1 AI Agent Identity.....	36
8.2 AI Permission Classes	37
8.3 Attribution Requirements.....	38
8.4 Prohibited Uses	39

9. Trust Framework	40
9.1 The Trust Score	41
9.2 Trust Levels	42
9.3 Verification Framework	43
9.4 Reputation Framework (Planned)	44
9.5 Regulatory Alignment and Trust Chain	45
10. APIs	46
10.1 Overview	47
10.2 Authentication	48
10.3 Core Resources	49
10.4 Status Codes	50
10.5 Events and Webhooks	51
11. Roadmap	52
11.1 Release Roadmap	53
11.2 Versioning Policy	54
11.3 Contributing	55
12. References	56

1. Vision

1.1 Vision Statement

HUGS ID™ exists to create a globally interoperable identity and governance standard that allows every individual to manage the lifecycle of their verified digital identity, while enabling trusted collaboration between institutions and artificial intelligence systems.

Mission: To become the trusted identity, provenance, and digital trust layer for the AI era by connecting verified people with verified digital content through a secure, transparent, interoperable, and auditable trust chain.

Every person already has a digital identity. Very few people control it. HUGS ID separates identity from data ownership: rather than creating another database, it establishes a governance layer that determines how identity, permissions, provenance, consent, verification, and trust interact across independent systems.

1.2 Objectives

- G-1 Define a canonical, machine-readable identity object that any institution or AI system can resolve, verify, and reference unambiguously.
- G-2 Give individuals enforceable authority over how verified information about them is used, by whom, and for what purpose.
- G-3 Make every AI action attributable to an identity, an authorization, a purpose, and a governance policy.
- G-4 Integrate with existing institutional systems — sports federations, universities, government agencies, healthcare providers, employers, research organizations, and AI platforms — rather than replace them.
- G-5 Evolve as an open technical standard that governments, institutions, and software providers can implement independently.

1.3 Non-Goals

- HUGS ID is not a centralized database of personal records. Source data remains with issuing institutions; the framework governs references, proofs, and permissions.
- HUGS ID is not a blockchain or cryptocurrency project. Cryptographic proofs are used for integrity and provenance, not for token economics.
- HUGS ID does not replace legal identity documents, credentials, or institutional records; it defines how those artifacts are verified, referenced, and governed.

2. Problem Statement

2.1 The Fragmentation of Digital Identity

Identity on today's Internet is stored inside thousands of disconnected databases owned by governments, schools, hospitals, employers, sports organizations, financial institutions, and technology companies. Each system holds a partial, unverifiable-from-outside view of a person, and none provides the individual with meaningful authority over how that information flows onward.

2.2 The AI Inflection Point

Artificial intelligence has dramatically accelerated the use of personal information. AI systems can summarize, classify, predict, recommend, and generate new content using data that individuals often neither control nor understand. Traditional identity systems were not designed to answer the questions the AI era now asks:

- Who may use this identity or record, under what authority, and for what purpose?
- Was this artifact created by a human, an institution, or an AI — and can that be proven?
- Which consent covered this specific use, and is it still in force?
- How much trust should a downstream system place in this claim?

HUGS ID addresses these questions by making identity, consent, provenance, and trust first-class, machine-readable objects with defined semantics.

2.3 Design Principles

Identity First — Every digital interaction begins with a verifiable identity. No governed action occurs against an anonymous or unresolved subject.

Data Sovereignty — Individuals retain authority over how verified information about them is governed. Institutional custody of a record does not imply unlimited rights to use it.

Verifiable Provenance — Every trusted record identifies who created it, when it was created, how it was verified, who approved it, and whether it has changed.

AI Accountability — Artificial intelligence must never operate without clearly defined permissions and governance. Every AI action is attributable to an identity, an authorization, a purpose, and a governance policy.

Interoperability — The framework integrates with existing systems rather than replacing them, and is designed for implementation by independent parties.

3. Architecture

3.1 Layered Architecture

The HUGS ID platform consists of independent layers. Each layer performs a specific function, exposes defined interfaces to the layers above and below it, and can be implemented or replaced without invalidating the rest of the stack. Requests flow upward through the stack; verification and trust flow downward.




Figure 3-1: The seven-layer HUGS ID architecture. The Identity Layer (teal) anchors the stack.

3.2 Layer Responsibilities

Layer	Responsibility	Primary Objects
Evidence Layer	Source records, provenance metadata, and integrity proofs anchoring all claims.	Evidence Record
Identity Layer	Canonical identity objects, identifier resolution, and lifecycle management.	Identity Object
Verification Layer	Validation of evidence and issuance of verifiable credentials and attestations.	Verification Object, Credential
Consent Layer	Machine-readable permissions granted by identity controllers, with purpose and revocation semantics.	Consent Object
Governance Layer	Policies, authority chains, roles, and accountability	Governance Object, Policy

Layer	Responsibility	Primary Objects
	for every governed action.	
AI Services	Models, agents, and automated decision systems operating under AI permission classes.	AI Permission Grant, Event Log
Applications	Federation, institutional, and consumer-facing systems consuming the layers below.	—

Table 3-1: Layer responsibilities and primary objects

3.3 Data Flow

1. An application or AI service requests a governed operation against an identity (read, derive, train, generate, distribute, or decide).
2. The Governance Layer evaluates the request against applicable policies and authority chains.
3. The Consent Layer confirms that a valid, unexpired consent covers the requested purpose and action class.
4. The Verification Layer confirms that the claims involved meet the required trust level.
5. The operation executes, and an immutable entry is written to the Event Log binding the action to its identity, authorization, purpose, and policy.

3.4 Governed Domains

The framework defines governed domains — vertical specializations that apply the core layers to a field of practice. Version 1.0 specifies Athlete Intelligence in detail (Section 7). Education, healthcare, employment, research, creative works, digital assets, and AI agents are recognized domains; additional domains may be added through future specifications.

4. Identity Layer

4.1 The Canonical Identity Object

DEFINITION — A Canonical Identity Object is the single, authoritative representation of a subject within the HUGS ID framework. It contains no personal records itself; it anchors references, verifications, consents, and governance decisions to one resolvable identifier.

Each Canonical Identity Object **MUST** contain the following fields:

Field	Type	Req.	Description
id	URI	MUST	Globally unique HUGS identifier (Section 4.2). Immutable once issued.
type	enum	MUST	individual organization ai-agent digital-asset
controller	URI	MUST	Identifier of the party exercising governance authority over the identity.
status	enum	MUST	Lifecycle state (Section 4.4): pending active suspended revoked retired.
created	date-time	MUST	ISO 8601 UTC timestamp of object creation.
updated	date-time	MUST	Timestamp of the most recent state change.
schema_version	string	MUST	Semantic version of the identity schema the object conforms to.
proof	object	SHOULD	Integrity proof over the object (hash or signature), enabling tamper evidence.

Table 4-1: Canonical Identity Object fields

Personal attributes (name, date of birth, nationality, and similar) are deliberately excluded from the Canonical Identity Object. They remain in institutional source systems and are referenced through verified claims, keeping the identity layer minimal, portable, and privacy-preserving.

4.2 Identifier Scheme

HUGS identifiers are URIs of the form:

`hugs:id:<type>:<namespace>:<unique-id>`

Example: `hugs:id:individual:us:8f3a9c2e-41d7`. The namespace segment is assigned by the issuing authority within a federation; the unique-id **MUST** be collision-resistant and **MUST NOT** itself encode personal information. Identifiers resolve through the Identity API (Section 10.3) to the current Canonical Identity Object and its status.

4.3 Identity Types

Type	Subject	Controller
individual	A natural person.	The person, or a legal guardian where applicable.
organization	An institution: federation, school, agency, employer, platform.	The institution's designated governance authority.
ai-agent	An AI system or agent acting under defined permissions (Section 8).	The operating organization; accountability is never waived.
digital-asset	A governed work: media, dataset, credential, creative output.	The rights holder, subject to applicable consents.

Table 4-2: Identity types and their controllers

4.4 Identity Lifecycle

State	Meaning	Permitted transitions
pending	Created, awaiting initial verification.	→ active, → revoked
active	Verified and resolvable; governed operations permitted.	→ suspended, → retired, → revoked
suspended	Temporarily inoperative pending review or dispute.	→ active, → revoked
revoked	Permanently invalidated (e.g., duplicate or fraud). Terminal.	—
retired	Voluntarily closed by the controller. Historical references remain resolvable.	—

Table 4-3: Lifecycle states and permitted transitions

REQ-4-1: An identity in any state other than active **MUST NOT** be the subject of a new governed operation. REQ-4-2: State transitions **MUST** be recorded in the Event Log with the acting identity and authority chain.

5. Governance Layer

5.1 Governance Objects

The Governance Layer defines the objects through which authority is expressed and enforced:

- Policy — a machine-readable rule set constraining operations within a domain or institution.
- Authority Chain — an ordered record of delegations from a root authority to the acting party.
- Role — a named capability bundle assignable to identities within an institution.
- Decision Record — the auditable outcome of a governance evaluation, stored in the Event Log.

5.2 Policy Model

A Policy binds a scope (domain, institution, or identity set) to constraints. Each Policy MUST define:

Field	Description
policy_id	Unique identifier and semantic version of the policy.
issuer	The authority chain of the issuing body.
scope	The identities, data classes, and domains to which the policy applies.
rules	Permitted and prohibited action classes, with conditions (e.g., minimum trust level).
effective / expires	Validity window. Expired policies are retained for audit but not enforced.
supersedes	Optional reference to a prior policy version this policy replaces.

Table 5-1: Policy object fields

Where multiple policies apply, the evaluation order is: individual controller policy, then institutional policy, then domain policy, then framework baseline. More specific policies override broader ones; prohibitions always override permissions.

5.3 Authority Chains

DEFINITION — An Authority Chain is a verifiable, ordered sequence of delegations — for example: framework → federation → club → registrar — establishing that the acting party's power to perform an operation derives from a recognized root authority.

REQ-5-1: Every governed operation MUST be traceable to a complete authority chain. REQ-5-2: A delegation cannot exceed the authority of its parent link; attenuation is permitted, amplification is not. REQ-5-3: Revoking any link suspends all downstream authority pending re-evaluation.

5.4 Institutional Roles

Role	Typical capabilities
Registrar	Create and retire identities within the institution's namespace.
Verifier	Validate evidence and issue attestations within a defined domain.
Governance Officer	Issue and amend policies; adjudicate disputes.
Data Custodian	Operate source systems referenced by evidence records.
Auditor	Read-only access to Event Logs and decision records.

Table 5-2: Reference institutional roles (extensible by policy)

6. Consent Layer

6.1 The Consent Object

DEFINITION — A Consent Object is a machine-readable record of a permission granted by the controller of an identity to a specified grantee, for a specified purpose, over a specified scope of verified information, for a defined period.

Field	Req.	Description
consent_id	MUST	Unique identifier of the consent.
subject	MUST	The identity to which the consented information pertains.
grantor	MUST	The identity (controller or authorized delegate) granting the consent.
grantee	MUST	The identity (institution, application, or AI agent) receiving the permission.
purpose	MUST	A term from the Purpose Taxonomy (Section 6.2).
scope	MUST	The data classes or specific claims covered by the consent.
actions	MUST	Permitted action classes (Section 8.2), e.g., read, derive, generate.
conditions	SHOULD	Additional constraints: minimum trust level, retention limits, redistribution terms.
granted / expires	MUST	Validity window. Expired consents are inoperative.
revoked	MUST	Revocation timestamp and reason, where applicable.

Table 6-1: Consent Object fields

6.2 Purpose Taxonomy

Consent purposes are drawn from a controlled vocabulary so that both humans and machines interpret permissions identically. The v1.0 taxonomy includes:

- eligibility-verification — confirming status, age, registration, or qualification.
- performance-analysis — computing metrics and insights over verified records.
- selection-evaluation — informing recruitment, admission, or selection decisions.
- media-licensing — use of identity-linked media or likeness under license.
- research — inclusion in anonymized or pseudonymized research datasets.
- model-training — use of verified information in AI training or fine-tuning (Section 8).

Domain specifications MAY extend the taxonomy; extensions MUST namespace custom purposes (e.g., athletics:transfer-clearance).

6.3 Consent Lifecycle

A consent moves through the states granted → active → (expired | revoked | superseded). A grantee **MUST** verify consent state at the time of each use; caching a consent evaluation beyond its stated conditions is non-conformant.

6.4 Revocation

Revocation is a first-class operation, not an exceptional one:

6. The controller revokes a consent through the Consent API; the object records the timestamp and reason.
7. The revocation propagates to the Event Log and to all registered webhooks of affected grantees (Section 10.5).
8. Downstream systems **MUST** cease new uses immediately and honor retention conditions for already-derived artifacts, as defined by the consent's conditions and applicable policy.

REQ-6-1: Consent **MUST** be as easy to withdraw as it was to grant. REQ-6-2: A use that continues after effective revocation **MUST** be detectable from the Event Log and attributable to the acting identity.

7. Athlete Intelligence

7.1 Overview

Athlete Intelligence is the first governed domain specified under the HUGS ID framework. It applies the Identity, Verification, Consent, Governance, and Trust layers to sport — a field where identity, age, eligibility, performance history, and media rights carry high value and are routinely mishandled across clubs, schools, federations, agents, and platforms.

The domain gives an athlete a single canonical identity that persists across teams, schools, leagues, and countries, while every claim attached to that identity — a registration, a result, a clearance, a highlight — carries verifiable provenance and a governing consent.

7.2 Governed Data Classes

Data class	Contents	Typical issuers
identity-registration	Verified identity, age band, registration status, federation membership.	Federations, clubs, schools
eligibility	Clearances, transfer status, amateur/professional standing, disciplinary records.	Federations, governing bodies
performance	Verified results, statistics, rankings, and milestones with provenance.	Event organizers, leagues, timing providers
development	Training benchmarks, combine results, academic standing where linked to eligibility.	Academies, schools, testing bodies
media-rights	Identity-linked media, likeness usage, licensing terms.	Athlete (controller), media partners

Table 7-1: Athlete Intelligence governed data classes

Health-adjacent information is out of scope for v1.0 and will be addressed in coordination with the healthcare domain specification.

7.3 Use Cases

- Eligibility verification — a federation confirms age band and registration in seconds against verified claims rather than documents of unknown provenance.
- Transfer and clearance — an authority chain from federation to club to registrar makes a clearance decision auditable end to end.
- Recruitment evaluation — scouts and programs consume verified performance histories under a selection-evaluation consent granted by the athlete.

- Media and likeness licensing — broadcasters and platforms operate under media-licensing consents with explicit revocation semantics.
- AI-assisted scouting — models access athlete data only under AI permission classes (Section 8), with every inference attributable and logged.

8. AI Permissions

8.1 AI Agent Identity

Every AI system that reads, derives from, trains on, generates about, distributes, or decides over governed information **MUST** hold an identity of type ai-agent (Section 4.3). An AI agent's controller is the operating organization; accountability for an agent's actions always resolves to a legal person or institution. Anonymous or unattributed AI operation against governed data is non-conformant.

8.2 AI Permission Classes

Class	Operation	Baseline requirement
read	Access verified data as context for inference.	Valid consent covering the purpose; active identity.
derive	Produce summaries, scores, or insights from verified data.	read + derive action in consent; output inherits provenance.
train	Use verified data in model training or fine-tuning.	Explicit model-training purpose; retention conditions honored.
generate	Create new content that references or depicts an identity.	generate action in consent; synthetic output MUST be labeled.
distribute	Publish or transmit derived or generated artifacts.	Distribution terms in consent conditions; downstream labeling preserved.
decide	Take an automated action affecting the subject (selection, eligibility, pricing).	Explicit policy authorizing the decision class; minimum trust level T3; human review path defined.

Table 8-1: AI permission classes and baseline requirements

8.3 Attribution Requirements

Every AI action within the framework **MUST** be attributable to four elements, recorded in the Event Log at execution time:

9. an identity — the ai-agent identity performing the action;
10. an authorization — the consent and permission classes relied upon;
11. a purpose — the taxonomy term the action serves;
12. a governance policy — the policy and authority chain under which the agent operates.

REQ-8-1: AI-generated content that references an identity **MUST** carry machine-readable provenance identifying it as generated, its source claims, and its permission basis. REQ-8-2: Permission evaluations **MUST** occur at the time of use; pre-computed or blanket authorizations are non-conformant.

8.4 Prohibited Uses

- Training on identity-linked data without an explicit model-training consent.
- Generating synthetic media of an identifiable person without a generate-class consent and labeling.
- Automated decisions in the decide class without an authorizing policy, T3 trust minimum, and a defined human review path.
- Circumventing revocation by retaining or re-deriving artifacts whose retention conditions have lapsed.
- Operating an AI agent without a resolvable identity and accountable controller.

9. Trust Framework

9.1 The Trust Score

DEFINITION — A Trust Score is a computed assessment of how much confidence a relying party should place in a claim, an identity, or an artifact. It is evidence-based, purpose-relative, and recomputed as underlying evidence, verification, and dispute state change.

A Trust Score is derived from five components:

- Verification depth — how many independent, authoritative sources confirm the claim.
- Provenance completeness — whether the full provenance record (creator, timestamp, verification method, approver, change history) is intact.
- Authority weight — the standing of the attesting institution within the relevant authority chain.
- Recency — the age of the evidence relative to the claim's expected validity window.
- Dispute state — whether the claim is uncontested, under review, or previously corrected.

9.2 Trust Levels

Level	Name	Basis
T0	Unverified	Self-asserted only; no supporting evidence. Usable for display, never for decisions.
T1	Basic	Single-source verification from a recognized custodian.
T2	Verified	Multi-source or institutional verification with complete provenance.
T3	Certified	Attested by the governing body through a complete authority chain. Required for decide-class AI actions.
T4	Canonical	Continuously verified, full provenance, uncontested over time. The reference state for high-stakes operations.

Table 9-1: Trust levels

Policies declare minimum trust levels per action class; the Governance Layer enforces them at evaluation time. A score is always scoped to a claim and a purpose — there is no universal “trusted person” flag in the framework.

9.3 Verification Framework

Verification converts evidence into attestations. The Verification Layer consumes Evidence Records — source documents, registry entries, event feeds — and issues Verification Objects recording what was

verified, by whom, against which method, and when. Verification Objects are the inputs from which trust levels are computed; they are versioned and revocable, and revocation of an underlying verification **MUST** trigger recomputation of dependent scores.

9.4 Reputation Framework (Planned)

Version 2.0 will introduce a Reputation Framework: longitudinal, behavior-based standing for institutions, verifiers, and AI agents — distinct from the per-claim Trust Score. Reputation models are explicitly deferred from v1.0 to keep the trust framework evidence-based and auditable.

9.5 Regulatory Alignment and Trust Chain

HUGS ID is designed around a fundamental principle:

Trust in artificial intelligence depends on trusted identity, verifiable provenance, and an unbroken chain of custody.

As governments and regulators introduce requirements for AI-generated content to include provenance metadata and verification mechanisms, HUGS ID provides a complementary trust framework. While provenance can indicate how digital content was created, HUGS ID extends that capability by helping establish who created it, who verified it, how it has changed over time, and whether it remains authentic.

The HUGS ID Trust Chain is built on five core pillars:

- **Verified Identity** — Establish the trusted identity of the creator, organization, athlete, or institution.
- **Verified Provenance** — Preserve standardized provenance metadata describing how, when, and with what tools digital content was created, including AI-assisted generation where applicable.
- **Chain of Custody** — Maintain an auditable record of significant events throughout the lifecycle of the asset, including creation, edits, approvals, transfers, publication, and revocation.
- **Integrity Verification** — Enable any recipient to independently verify that digital content remains consistent with its documented history and has not been altered outside the recorded chain of custody.
- **Transparent Attribution** — Provide clear, verifiable information regarding authorship, ownership, licensing, permissions, and AI involvement.

Together, these five pillars establish a Digital Trust Chain that extends beyond watermarking or embedded metadata alone. Rather than simply indicating that content was generated by AI, HUGS ID links digital assets to verified people, organizations, permissions, and lifecycle events. This creates a transparent and auditable framework for authenticity, accountability, and trust across sports, entertainment, education, government, and enterprise ecosystems.

The platform is intended to evolve alongside emerging standards for AI transparency, provenance, and responsible use while remaining interoperable with future regulatory and industry frameworks.

This formulation establishes an architectural foundation for HUGS ID — shifting the industry narrative from passive AI detection to proactive, lifecycle-wide accountability.

10. APIs

10.1 Overview

HUGS ID exposes a versioned REST API over HTTPS. All resources are namespaced under `/v1/`; object representations are JSON and conform to the published v1 schemas. API behavior follows the same governance rules as the rest of the framework — the API is an enforcement surface, not a bypass.

10.2 Authentication

Machine clients authenticate with OAuth 2.0 client credentials; user-facing flows use OpenID Connect. Tokens carry the caller's identity and authority context, which are evaluated on every request. REQ-10-1: API credentials are issued per identity and role; shared or anonymous credentials are non-conformant.

10.3 Core Resources

Resource	Methods	Purpose
<code>/v1/identities</code>	GET, POST	Resolve, search, and register Canonical Identity Objects.
<code>/v1/identities/{id}</code>	GET	Retrieve the current object and lifecycle state of an identity.
<code>/v1/identities/{id}/verifications</code>	GET	List verifications and attestations attached to an identity.
<code>/v1/consents</code>	GET, POST	Grant consents; list consents by subject, grantor, or grantee.
<code>/v1/consents/{id}/revoke</code>	POST	Revoke a consent with a reason; propagates to the Event Log and webhooks.
<code>/v1/governance/policies</code>	GET, POST	Publish and query policies within the caller's authority scope.
<code>/v1/trust/scores/{identityId}</code>	GET	Retrieve claim-level trust scores with component breakdown.
<code>/v1/events</code>	GET	Query the Event Log (audit stream) within the caller's authorization scope.

Table 10-1: Core API resources (v1)

10.4 Status Codes

Code	Meaning
200 OK	Request succeeded.

Code	Meaning
201 Created	Object created (identity registered, consent granted, policy published).
400 invalid_request	Malformed request or schema validation failure.
401 unauthorized	Missing or invalid credentials.
403 consent_denied	No valid consent covers the requested purpose and action class.
403 policy_denied	A governing policy prohibits the operation.
403 trust_insufficient	The claims involved do not meet the required trust level.
404 not_found	Identifier does not resolve.
409 version_conflict	Object has changed since the caller's version; re-read and retry.
410 revoked	Target object (identity, consent, verification) has been revoked.
429 rate_limited	Request quota exceeded.

Table 10-2: Standard status codes

10.5 Events and Webhooks

Every governed action emits an event to the Event Log. Subscribers register webhooks for event types including: `identity.created`, `identity.updated`, `identity.state_changed`, `consent.granted`, `consent.revoked`, `verification.issued`, `verification.revoked`, `trust.changed`, `policy.updated`, and `ai.action_executed`. Webhook payloads are signed so receivers can verify framework origin.

11. Roadmap

11.1 Release Roadmap

Release	Scope
Version 1.0 (this document)	Identity Framework; Governance Framework; Consent Framework; Verification Model; Athlete Intelligence domain; Trust Chain (five pillars).
Version 2.0	AI Agent Governance (full agent lifecycle); Cross-Domain Trust Exchange; Distributed Verification; Institutional Federation; Advanced Reputation Models.
Version 3.0	Global Identity Federation; Autonomous AI Governance; Real-Time Compliance Engine; Decentralized Trust Networks.

Table 11-1: Release roadmap

11.2 Versioning Policy

Every specification and schema follows semantic versioning (MAJOR.MINOR.PATCH). MAJOR releases may change object semantics; MINOR releases add backward-compatible capabilities; PATCH releases are editorial. Backward-compatibility requirements are defined for every release: within a major line, consumers written against v1.0 MUST continue to operate against all v1.x schemas without modification.

11.3 Contributing

HUGS ID welcomes contributions from developers, researchers, academic institutions, governing bodies, standards organizations, and industry partners. Independent implementations are encouraged, provided they conform to the published specifications. Future versions will evolve through public review and collaborative implementation; comments and implementation reports should be submitted through the documentation portal.

12. References

Normative References

- [1] S. Bradner, “Key words for use in RFCs to Indicate Requirement Levels,” BCP 14, RFC 2119, IETF, March 1997. <https://www.rfc-editor.org/rfc/rfc2119>

Informative References

- [2] W3C, “Verifiable Credentials Data Model v2.0,” W3C Recommendation. <https://www.w3.org/TR/vc-data-model-2.0/>
- [3] W3C, “Decentralized Identifiers (DIDs) v1.0,” W3C Recommendation. <https://www.w3.org/TR/did-core/>
- [4] D. Hardt (Ed.), “The OAuth 2.0 Authorization Framework,” RFC 6749, IETF, October 2012. <https://www.rfc-editor.org/rfc/rfc6749>
- [5] OpenID Foundation, “OpenID Connect Core 1.0.” https://openid.net/specs/openid-connect-core-1_0.html
- [6] ISO/IEC 27001, “Information security management systems — Requirements.”
- [7] NIST Special Publication 800-63, “Digital Identity Guidelines.”
- [8] Regulation (EU) 2016/679 (General Data Protection Regulation).
- [9] Coalition for Content Provenance and Authenticity (C2PA), “C2PA Technical Specification — Content Credentials.” <https://c2pa.org/specifications/>

These references inform the design of HUGS ID. Where this specification and a referenced standard conflict, this specification defines behavior within the HUGS ID framework only and does not modify the referenced standard.



hugs id

Technical Specification · Version 1.0 · Public Working Draft



hugsid.com · learn.xpguess.com

© 2026 HUGS ID™. This specification is published as a public working draft for review and implementation.

